

MegaPipe : A New Programming Interface for Scalable Network I / O

Sangjin Han⁺, Scott Marshall⁺, Byung-Gon Chun^{*}, and Sylvia Ratnasamy⁺

⁺University of California, Berkeley; ^{*}Yahoo! Research

OSDI, 2012

Presented By Dong Yuan & Zhihui Deng
(2015210938 2015210926)

Background

- Message-Oriented Workload
 - Short connections or small messages
 - Examples: HTTP, RPC, DB
- Issues with message-oriented workloads
 - System call overhead
 - Shared listening socket
 - File abstraction overhead

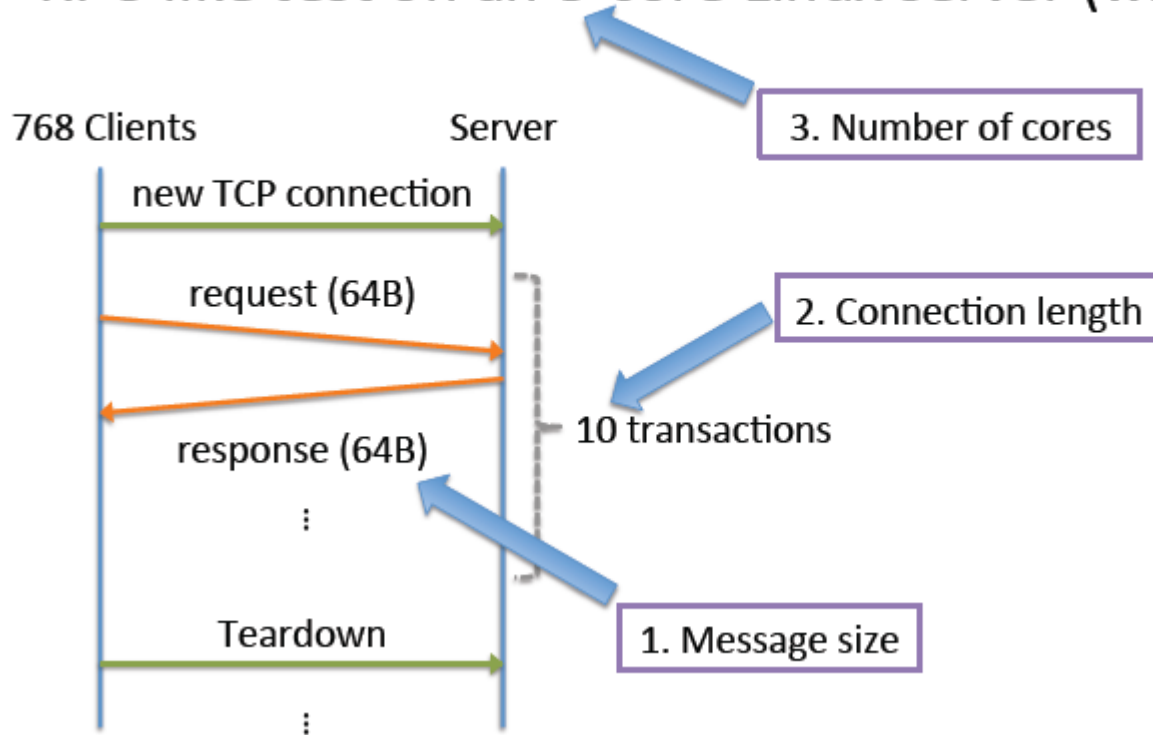


Solved Issues Comparison with mTCP & Fastsocket

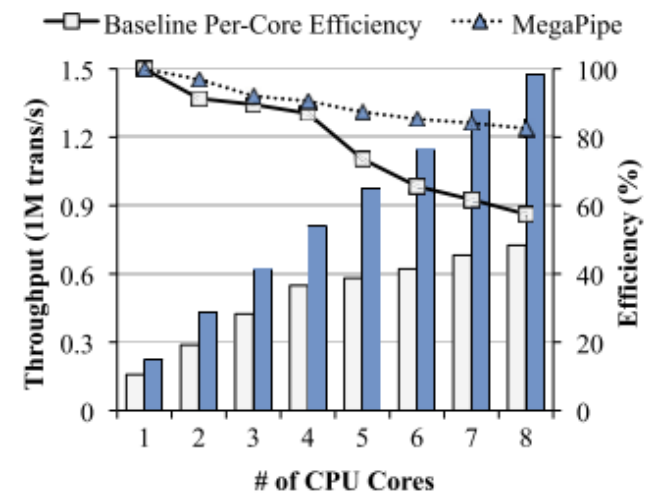


- MegaPipe (OSDI, 2012)
 - System call overhead
 - Shared listening socket
 - File abstraction overhead
- mTCP (NSDI, 2014)
 - Shared resources
 - Broken locality
 - Per packet processing
- Fastsocket (ASPLOS, 2016)
 - Shared resources
 - Broken locality
 - Uncompatible API

Microbenchmark

RPC-like test on an 8-core Linux server (with epoll)



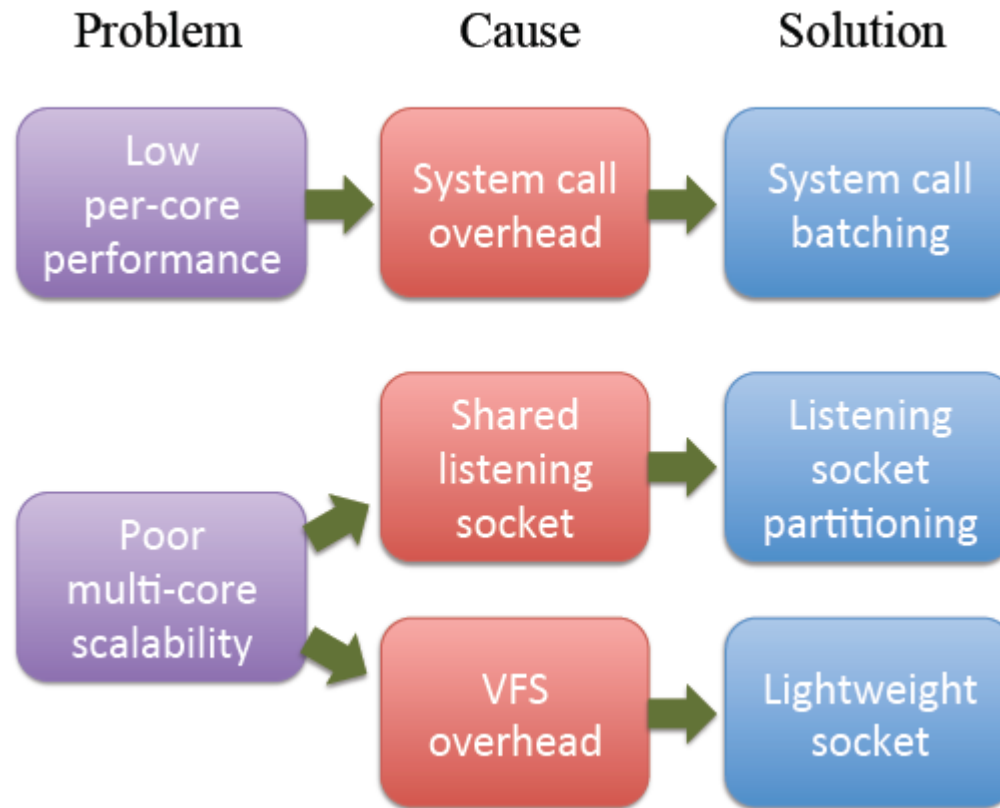
Performance of Message-Oriented Workloads



Design Goal

- API, applicable to existing event-driven server applications with moderate efforts
- Unified interface for various I/O types, TCP connection, UNIX domain sockets, disk files...
- Low overhead & multi-core scalability

Overview




Key Primitives

- ***Handle***

- Similar to file descriptor
 - But only valid within a channel
- TCP connection, pipe, disk file...

- ***Channel***

- A per-core, bi-directional pipe between the kernel and user
- Multiplexes I/O operations of its handles



Completion Notification Model

- Application issue asynchronous I/O commands
- Kernel notifies the application when the commands are complete
- Why CNM?
 - CNM allows transparent batching of I/O commands and notifications
 - It is compatible with not only sockets but also disk files
 - Simplify the complexity of I/O multiplexing

```
epoll_ctl(fd1, EPOLLIN);  
epoll_ctl(fd2, EPOLLIN);  
epoll_wait(...);
```

```
...  
ret1 = recv(fd1, ...);  
...  
ret2 = recv(fd2, ...);  
...
```



Readiness Model

```
mp_read(handle1, ...);  
mp_read(handle2, ...);
```



...

```
ev = mp_dispatch(channel);
```

...

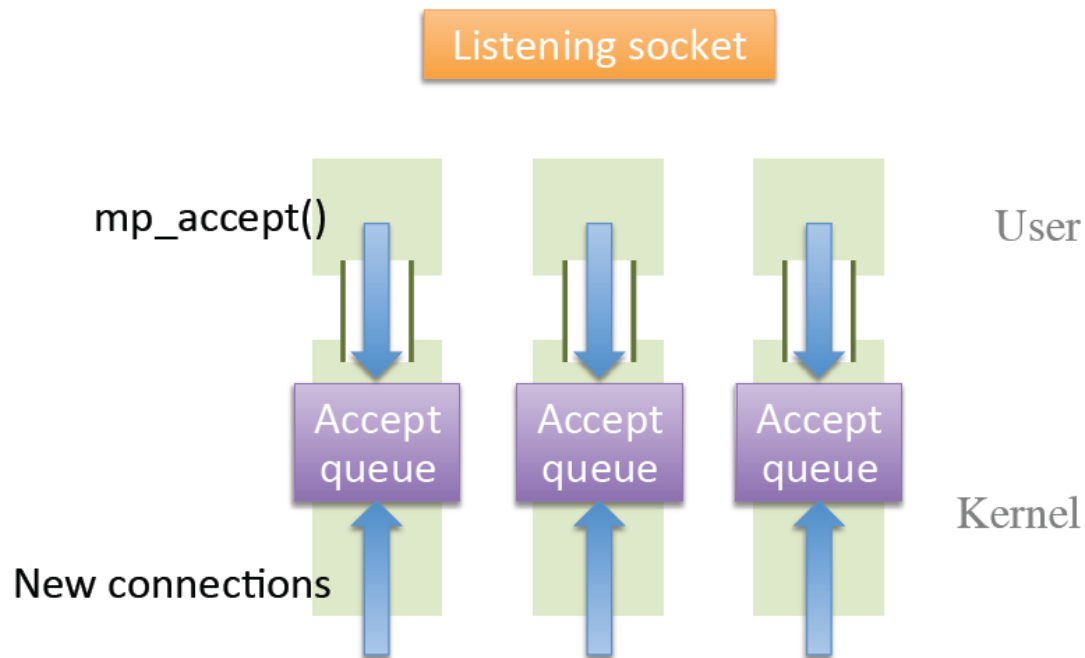
```
ev = mp_dispatch(channel);
```

...

CNM

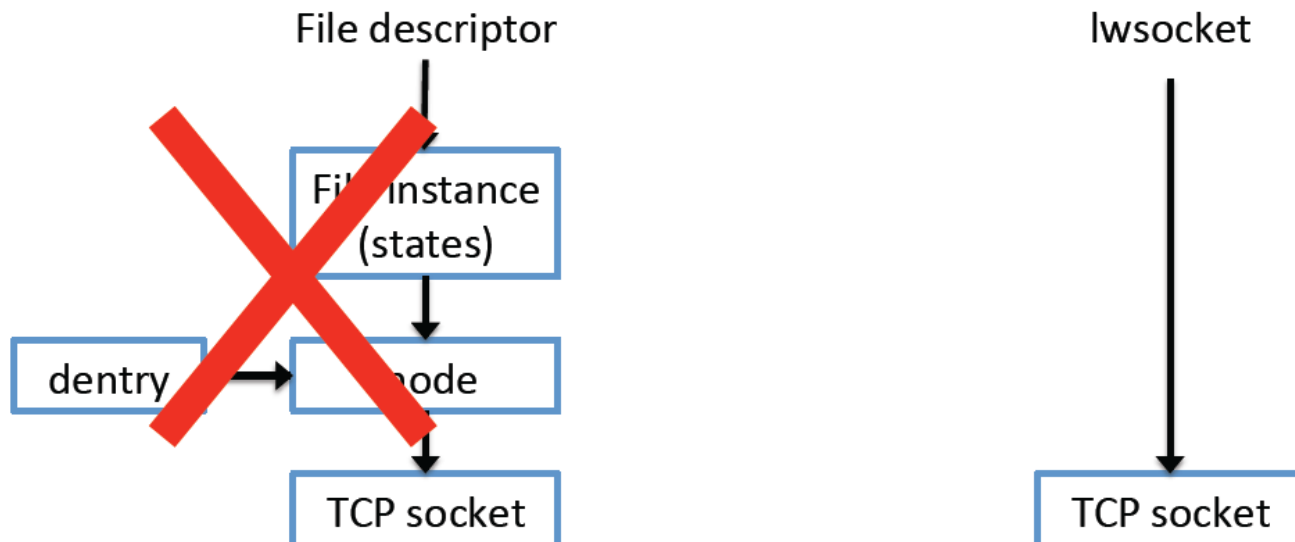
Listening Socket Partitioning

- Per-core accept queue for each channel
 - Instead of the globally shared accept queue



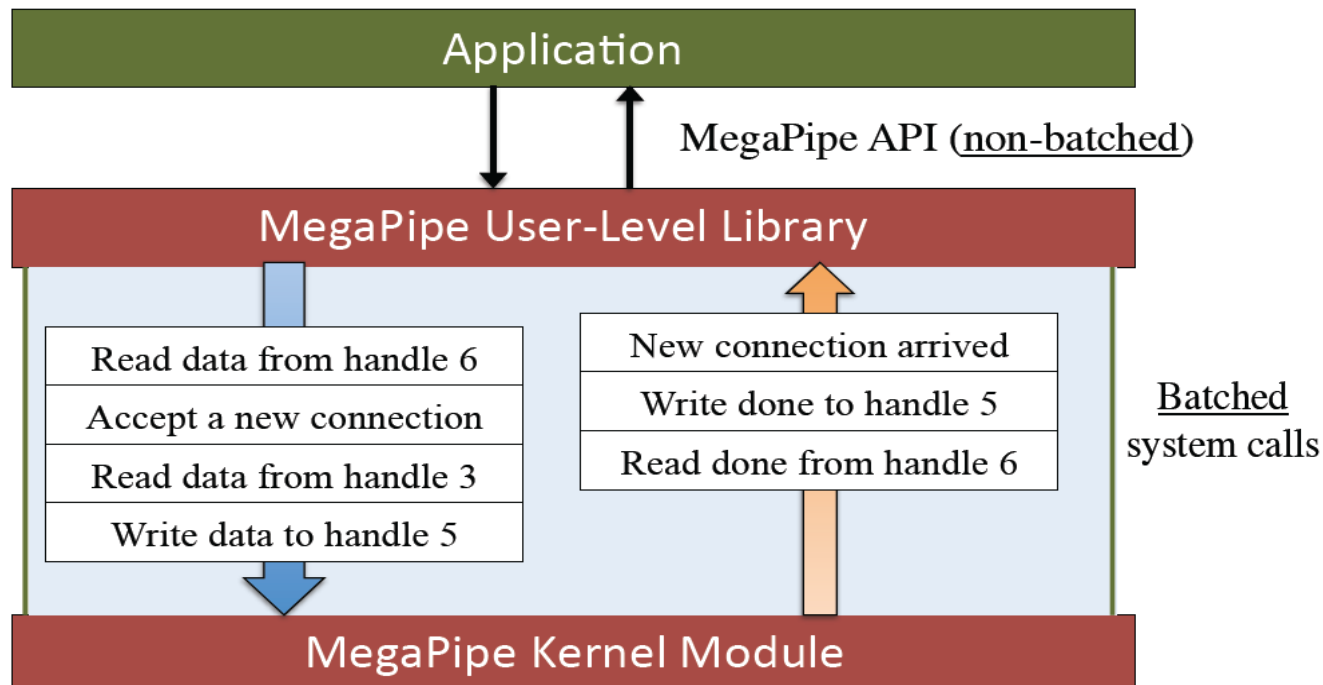
Lightweight Socket

- Sockets are ephemeral and rarely shared
 - Bypass the VFS layer
 - Convert into a regular file descriptor only when necessary



System Call Batching

- System calls are expensive due to cost of mode switching and bad cache locality
- Transparent batching




Implementation

- Kernel
 - One kernel module (~1800 lines)
 - Kernel itself (~400 lines)
- User-Level Library
 - ~400 lines
- Application
 - Supportive to event-driven server
 - ~hundreds of lines



Evaluation

- Multi-core scalability
 - Throughput improvement with various message sizes




Evaluation

- Multi-core scalability
 - Throughput improvement with various connection lengths (# of transactions)




Application Evaluation

- Memcached



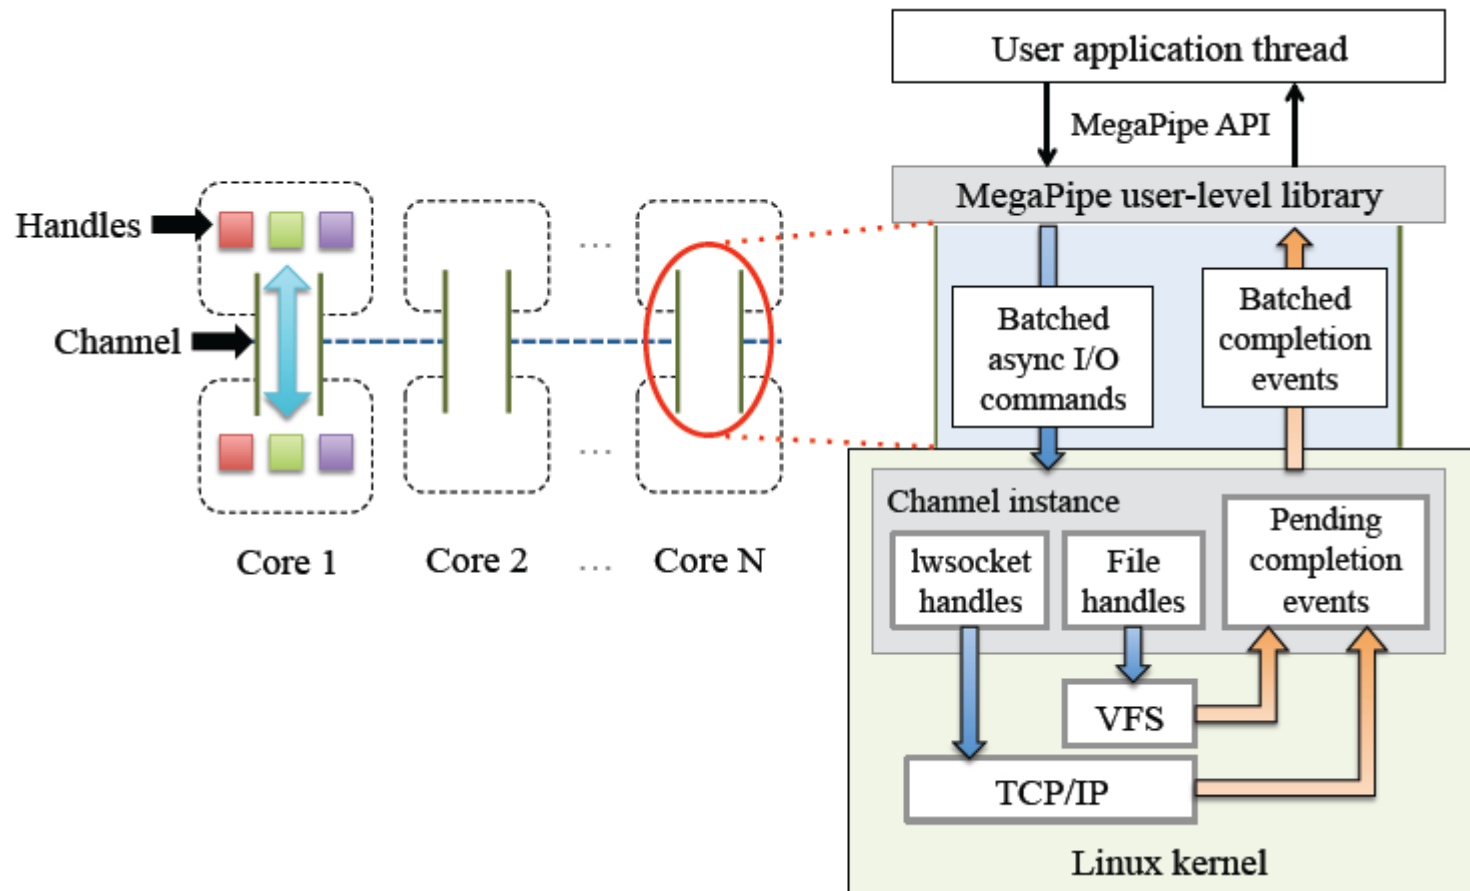
- nginx



Conclusion

- MegaPipe
 - Key abstraction: per-core channel
 - Enabling three optimization:
 - Batching, partitioning, lwsocket
 - Performance improvement in multi-core scalability and application

Comments About This Paper



Thank you!