

# MegaPipe : A New Programming Interface for Scalable Network I / O

Sangjin Han<sup>+</sup>, Scott Marshall<sup>+</sup>, Byung-Gon Chun<sup>\*</sup>, and Sylvia Ratnasamy<sup>+</sup>

<sup>+</sup>University of California, Berkeley; <sup>\*</sup>Yahoo! Research

OSDI, 2012

Presented By Dong Yuan, 2015210938

# Background

- Message-Oriented Workload
  - Short connections or small messages
    - Examples: HTTP, RPC, DB
- Issues with message-oriented workloads
  - System call overhead
  - Shared listening socket
  - File abstraction overhead

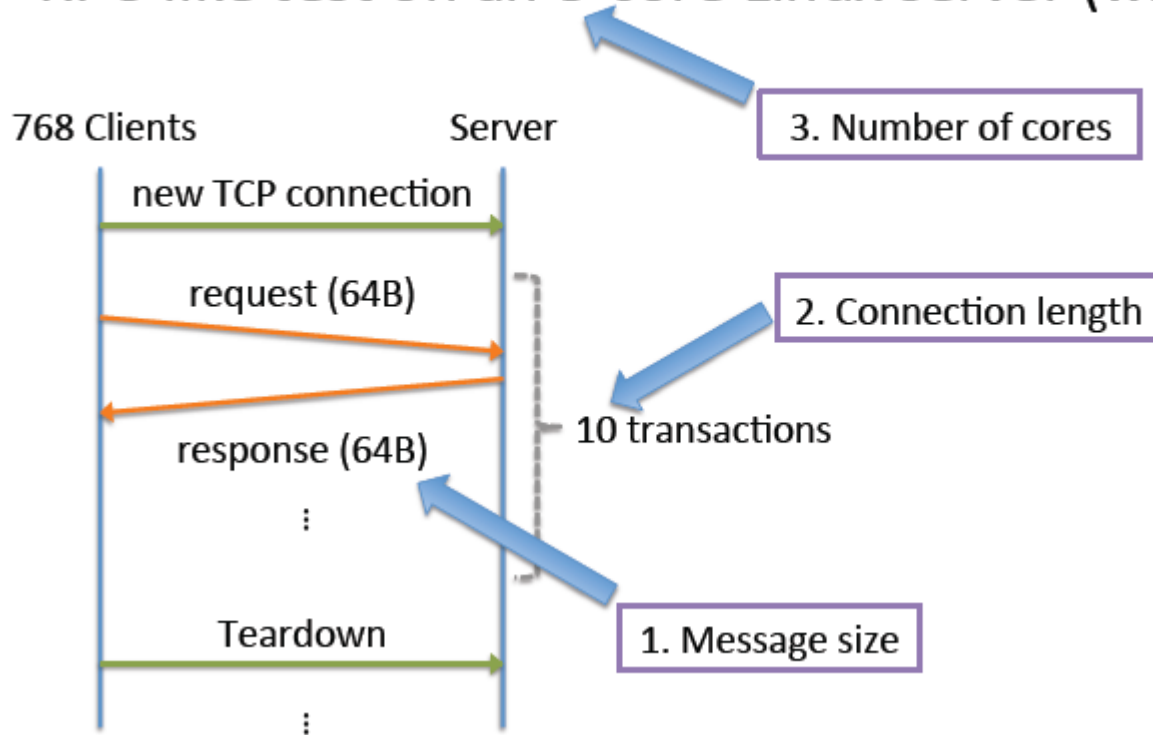


# Solved Issues Comparison with mTCP & Fastsocket

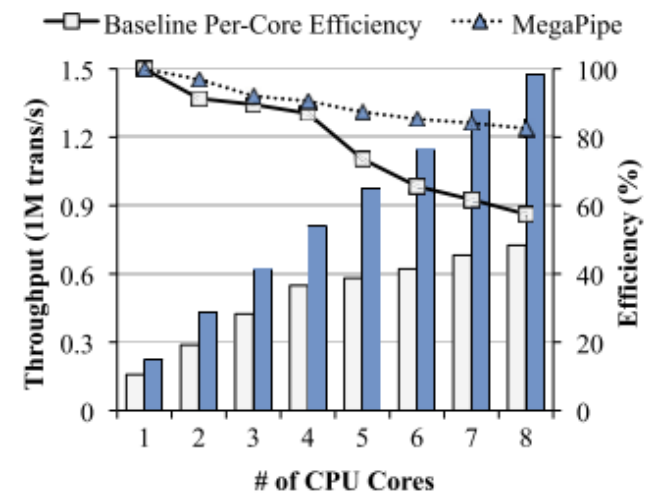


- MegaPipe (OSDI, 2012)
  - System call overhead
  - Shared listening socket
  - File abstraction overhead
- mTCP (NSDI, 2014)
  - Shared resources
  - Broken locality
  - Per packet processing
- Fastsocket (ASPLOS, 2016)
  - Shared resources
  - Broken locality
  - Uncompatible API

# Microbenchmark

## RPC-like test on an 8-core Linux server (with epoll)



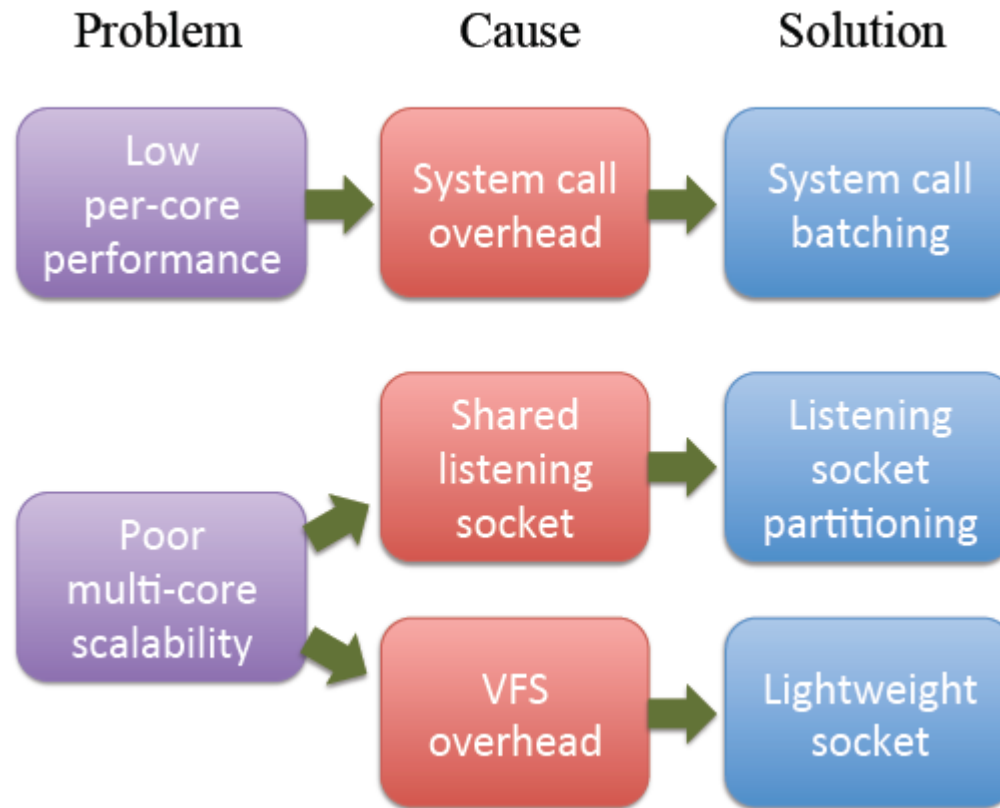
# Performance of Message-Oriented Workloads



# Design Goal

- API, applicable to existing event-driven server applications with moderate efforts
- Unified interface for various I/O types, TCP connection, UNIX domain sockets, disk files...
- Low overhead & multi-core scalability

# Overview




# Key Primitives

- **Handle**

- Similar to file descriptor
  - But only valid within a channel
- TCP connection, pipe, disk file...

- **Channel**

- A per-core, bi-directional pipe between the kernel and user
- Multiplexes I/O operations of its handles





# Completion Notification Model

- Application issue asynchronous I/O commands
- Kernel notifies the application when the commands are complete
- Why CNM?
  - CNM allows transparent batching of I/O commands and notifications
  - It is compatible with not only sockets but also disk files
  - Simplify the complexity of I/O multiplexing

```
epoll_ctl(fd1, EPOLLIN);  
epoll_ctl(fd2, EPOLLIN);  
epoll_wait(...);
```

```
...  
ret1 = recv(fd1, ...);  
...  
ret2 = recv(fd2, ...);  
...
```



Readiness Model

```
mp_read(handle1, ...);  
mp_read(handle2, ...);
```



...

```
ev = mp_dispatch(channel);
```

...

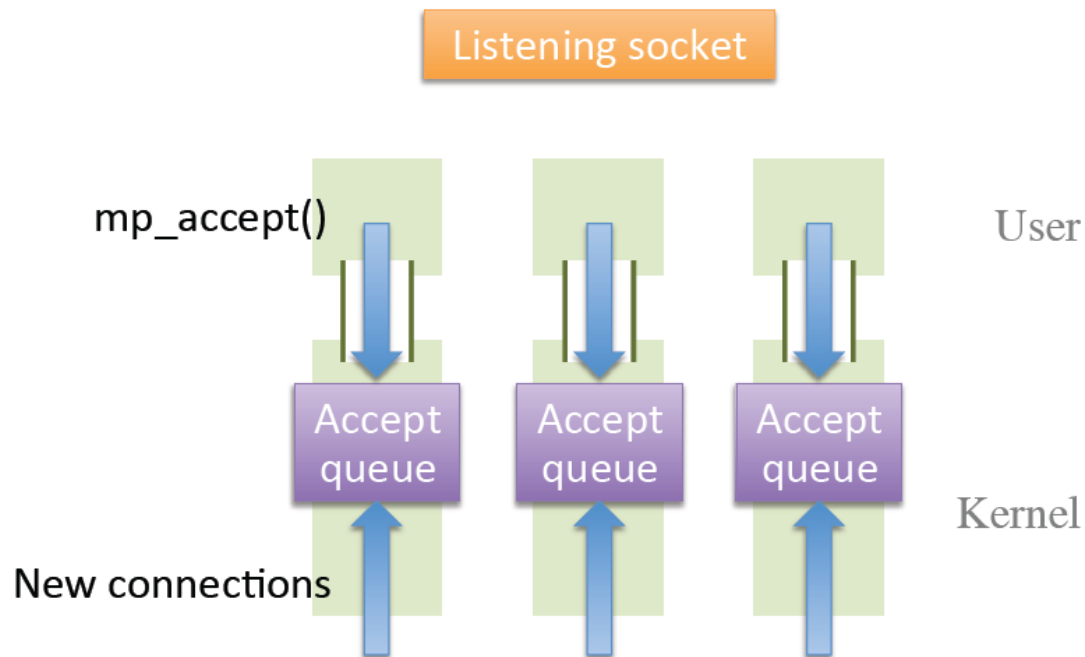
```
ev = mp_dispatch(channel);
```

...

CNM

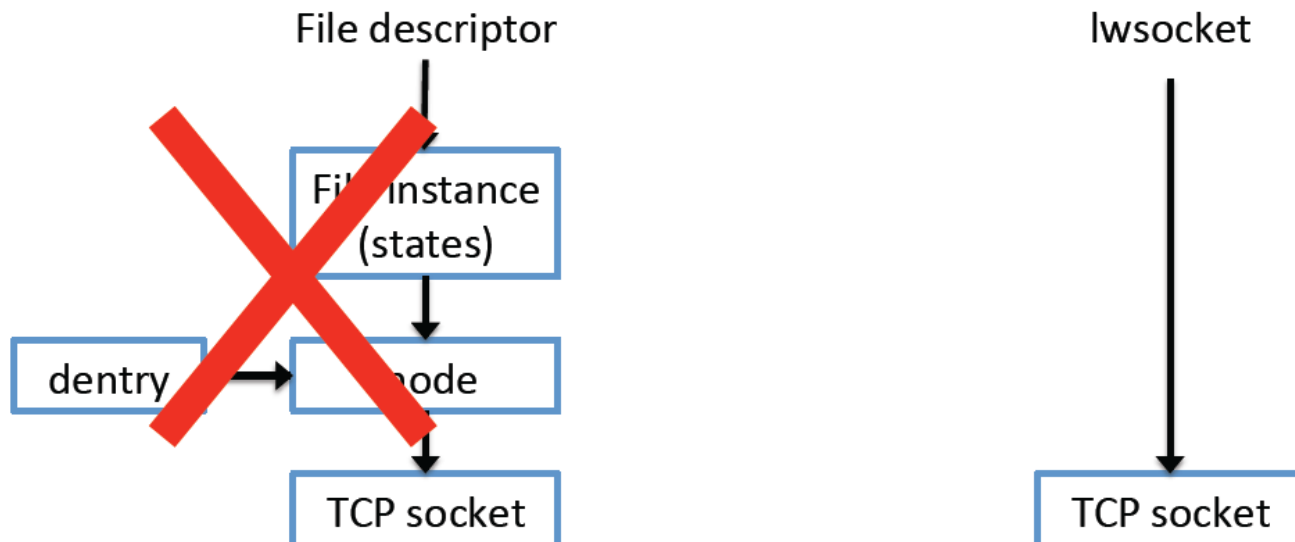
# Listening Socket Partitioning

- Per-core accept queue for each channel
  - Instead of the globally shared accept queue



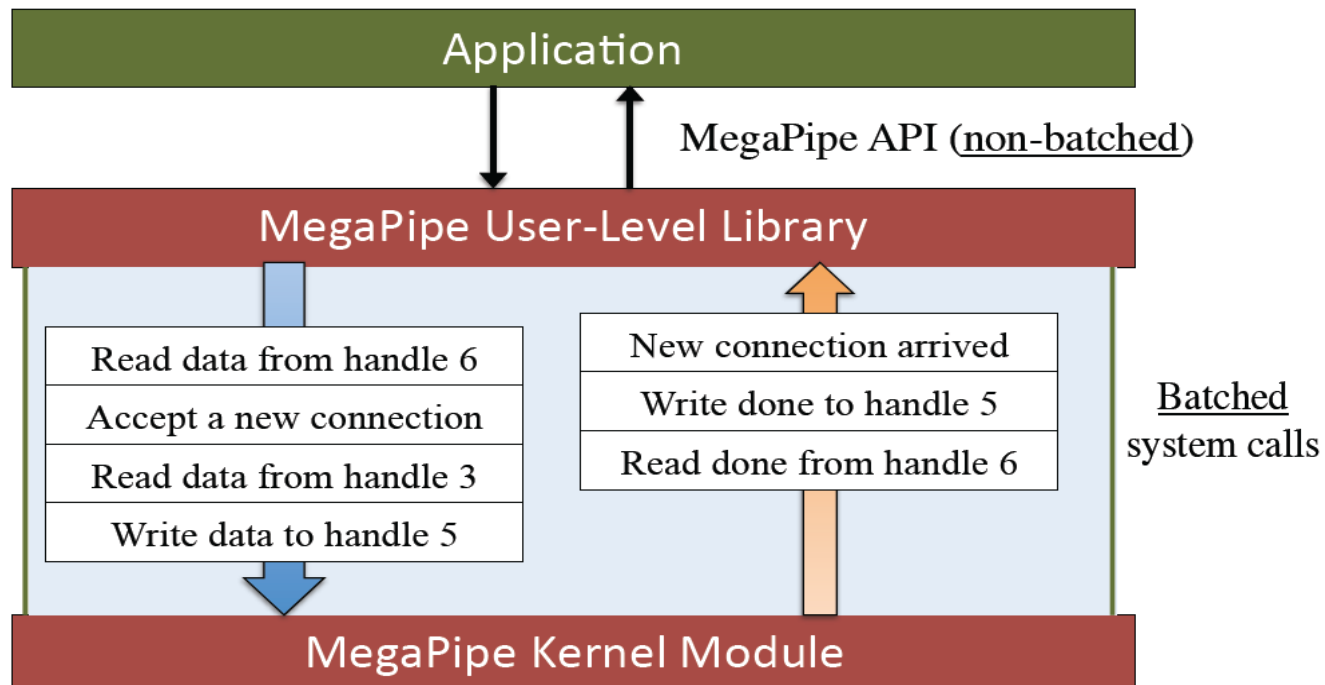
# Lightweight Socket

- Sockets are ephemeral and rarely shared
  - Bypass the VFS layer
  - Convert into a regular file descriptor only when necessary



# System Call Batching

- System calls are expensive due to cost of mode switching and bad cache locality
- Transparent batching




# Implementation

- Kernel
  - One kernel module (~1800 lines)
  - Kernel itself (~400 lines)
- User-Level Library
  - ~400 lines
- Application
  - Supportive to event-driven server
  - ~hundreds of lines


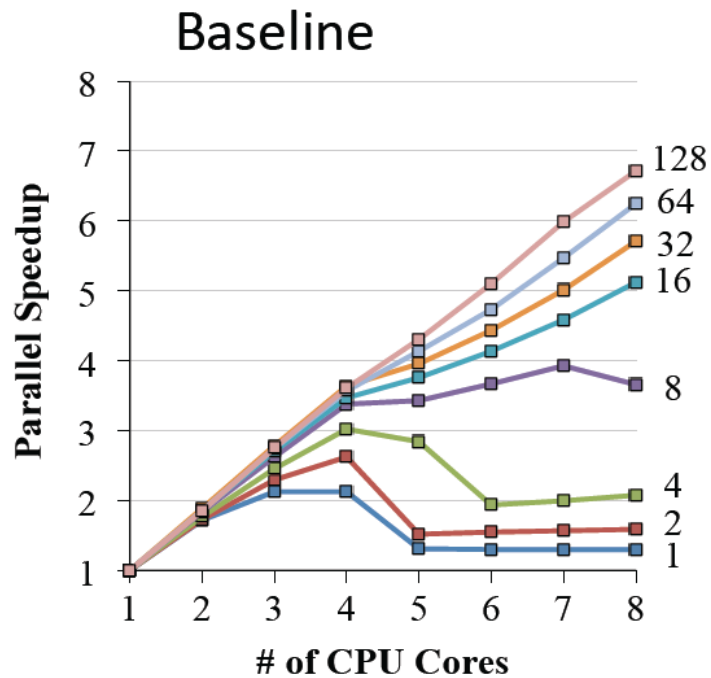
# Evaluation

- Multi-core scalability
  - Throughput improvement with various message sizes




# Evaluation

- Multi-core scalability
  - Throughput improvement with various connection lengths (# of transactions)




# Application Evaluation

- Memcached



- nginx





# Conclusion

- MegaPipe
  - Key abstraction: per-core channel
  - Enabling three optimization:
    - Batching, partitioning, lwsocket
  - Performance improvement in multi-core scalability and application

# Comments About This Paper

